Bias-Variance Tradeoff

Mathematics for Machine Learning

Lecturer: Matthew Wicker

Logistics

- Thank you to all who filled out the survey!
- Our final lecture will be an overview lecture where we connect all of the topics and processes together
- After the final lecture, GTA's will be assigned problem topics to discuss in small groups
- You will get the following materials before next Friday's lecture: Practice Exam, formula sheet, slides containing list of examinable material

Material Covered

Models: Linear models, basis expansion, logistic regression, neural networks, Prob. densities, Bayesian density estimation

Techniques: Least squares estimation, forward AD, reverse AD, computational graphs, gradient descent, convergence, convexity, Lipschitz continuity, Maximum likelihood, maximum a posteriori, Bayesian inference, LOTUS, change of variables, expectation identities, equating coefficients, joint Gaussian, epistemic/aleatoric uncertainty, concentration inequalities, cross-validation, regularization, generalization error bounds

Settings: Regression, Classification, Density Estimation

This lecture: Estimators, bias of an estimator, variance of an estimator, bias variance trade off

Recall from last lectures

it's better but worst generalisation

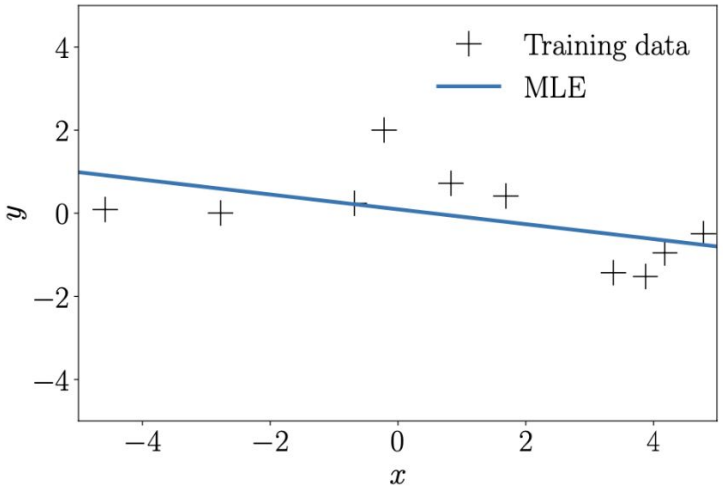

$\boldsymbol{\phi}(x) = [1\ x]^\intercal$

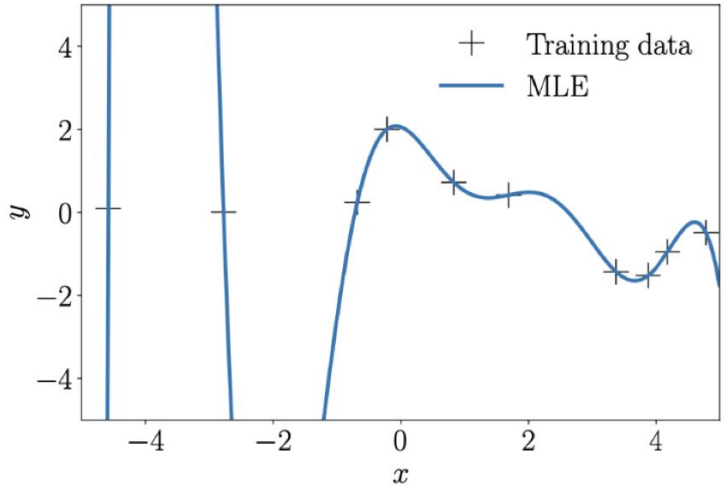

$\boldsymbol{\phi}(x) = [1\ x\ x^2\ x^3, \ldots]^\intercal$

Recall from last lectures

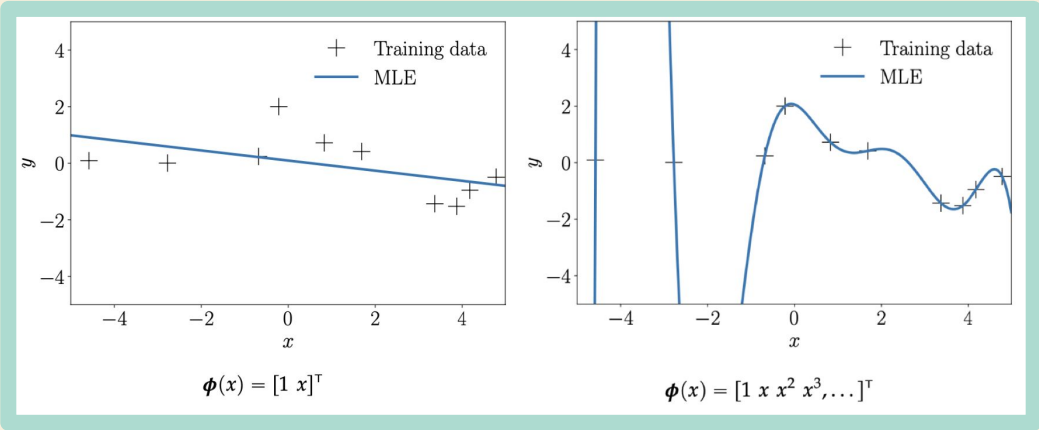

Model validation ideas:

- Test-set -> generalization error w/concentration
- Generalization error bounds for class of models
- Regularization (implicit/explicit)
- *Cross-validation*

We talk of estimators, let's define them

Definition 2.1. Statistic A statistic S is a random variable that is a function of some data $\mathcal{D}$, $S = g(\mathcal{D})$ where the data $\mathcal{D}$ is a collection of random variables.

An estimator is a statistic that aims at approximating a parameter/property of a distribution.

RV

↓

$S = g(D)$

estimator -> S that aims dist. param.

We talk of estimators, let's define them

Definition 2.1. Statistic A statistic S is a random variable that is a function of some data $\mathcal{D}$, $S = g(\mathcal{D})$ where the data $\mathcal{D}$ is a collection of random variables.

An estimator is a statistic that aims at approximating a parameter/property of a distribution.

sample mean estimator

$$\hat{Z} = \frac{Z_1 + Z_2 + \ldots + Z_N}{N}$$

$\int^n$ applied to data

We talk of estimators, let's define them

An estimator is a statistic that aims at approximating a parameter/property of a distribution.

estimator of Z

$$\text{Bias}(\hat{Z}_n) = \mathbb{E}[\hat{Z}_n - Z]$$

Expected diff between estimator and actual statistic

We talk of estimators, let's define them

An estimator is a statistic that aims at approximating a parameter/property of a distribution.

$$\text{Bias}(\hat{Z}_n) = \mathbb{E}[\hat{Z}_n - Z] = 0$$

Unbiased estimator

We proved unbiased estimator for the sample mean

From lecture 10:

$$\hat{Z} = \frac{Z_1 + Z_2 + \ldots + Z_N}{N}$$

$$\mathbb{E}[\hat{Z}] = \frac{\mathbb{E}[Z_1] + \mathbb{E}[Z_2] + \ldots + \mathbb{E}[Z_N]}{N} \quad \text{i.i.d}$$

$$= \frac{\mu_1 + \mu_2 + \ldots \mu_N}{N} = \mu$$

Sample variance estimators

$$\hat{\sigma^2}_n = \frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2 \qquad \hat{\sigma^2}'_n = \frac{1}{n-1}\sum_{i=1}^{n}(X_i - \mu)^2$$

Here we have two sample variance estimators, but which one is the biased estimator and which is unbiased?

Letting: $\mu = \hat{\mu}_n$

Sample variance estimators

Let's check the simpler looking equation for bias

$$\hat{\sigma^2}_n = \frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2$$

$$\text{bias}(\hat{\sigma^2}_n) = \mathbb{E}[\hat{\sigma^2}_n - \sigma^2]$$

Bias: $E[\hat{S} - S]$

$= E[\hat{S}] \cdot E[S] : E[\hat{S}] - S$

Var: $E[(\hat{S} - E[\hat{S}])^2]$

Err: $E[(\hat{S} - S)^2]$

$= \text{Var} + \text{Bias}^2$

$= E[(\hat{S} - E[\hat{S}])^2] + (E[\hat{S}] - S)^2$

Sample variance estimators

Let's check the simpler looking equation for bias

$$\hat{\sigma^2}_n = \frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2$$

$$\begin{aligned}\text{bias}(\hat{\sigma^2}_n) &= \mathbb{E}[\hat{\sigma^2}_n - \sigma^2] \\ &= \mathbb{E}[\frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2 - \sigma^2]\end{aligned}$$

Sample variance estimators

Let's check the simpler looking equation for bias

$$\text{bias}(\hat{\sigma^2}_n) = \mathbb{E}[\hat{\sigma^2}_n - \sigma^2]$$

$$= \mathbb{E}[\frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2 - \sigma^2]$$

$$= \mathbb{E}[\frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2] - \sigma^2$$

$E(\sigma^2) = \sigma^2$

Sample variance estimators

Simplifying the expression in the expectation:

$$\frac{1}{n}\sum_{i=1}^{n}(X_i-\mu)^2=\frac{1}{n}\sum_{i=1}^{n}(X_i^2-2X_i\mu+\mu^2)$$

Sample variance estimators

Simplifying the expression in the expectation:

$$\frac{1}{n}\sum_{i=1}^{n}(X_i-\mu)^2 = \frac{1}{n}\sum_{i=1}^{n}(X_i^2-2X_i\mu+\mu^2)$$

$$\propto \sum_{i=1}^{n}X_i^2 - \sum_{i=1}^{n}2X_i\mu + \sum_{i=1}^{n}\mu^2$$

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$\frac{1}{n}\sum_{i=1}^{n}(X_i-\mu)^2 = \frac{1}{n}\sum_{i=1}^{n}(X_i^2-2X_i\mu+\mu^2)$$

$$\propto \sum_{i=1}^{n}X_i^2 - \sum_{i=1}^{n}2X_i\mu + \sum_{i=1}^{n}\mu^2$$

$$= \left(\sum_{i=1}^{n}X_i^2\right) - 2\mu\left(\sum_{i=1}^{n}X_i\right) + n\mu^2$$

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$= \left(\sum_{i=1}^{n} X_i^2\right) - 2\mu\left(\sum_{i=1}^{n} X_i\right) + n\mu^2$$

$$= \left(\sum_{i=1}^{n} X_i^2\right) - 2n\mu^2 + n\mu^2$$

$\mu = \frac{\sum x_i}{n}$

$\hookrightarrow \sum x_i = n\mu$

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$= \left(\sum_{i=1}^{n} X_i^2\right) - 2\mu\left(\sum_{i=1}^{n} X_i\right) + n\mu^2$$

$$= \left(\sum_{i=1}^{n} X_i^2\right) - 2n\mu^2 + n\mu^2$$

$$= \left(\sum_{i=1}^{n} X_i^2\right) - n\mu^2$$

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$\mathbb{E}\left[\left(\sum_{i=1}^{n} X_i^2\right) - n\mu^2\right] = \left(\sum_{i=1}^{n} \mathbb{E}[X_i^2]\right) - \mathbb{E}[n\mu^2]$$

Expectation over data (X_i)

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$\mathbb{E}\left[\left(\sum_{i=1}^{n} X_i^2\right) - n\mu^2\right] = \left(\sum_{i=1}^{n} \mathbb{E}[X_i^2]\right) - \mathbb{E}[n\mu^2]$$

$\sigma^2 = E[X^2] - E[X]^2$

μ^2

Aside, identities we will prove later:

$$\mathbb{E}[X_i^2] = \sigma^2 + \mu^2$$

Hint: this is just rearranging the variance definition

$$\mathbb{E}[\mu^2] = \frac{\sigma^2}{n} + \mu^2$$

Sample variance estimators

Simplifying the expression in the expectation:

$$= \left(\sum_{i=1}^{n} \mathbb{E}[X_i^2] \right) - \mathbb{E}[n\mu^2]$$

$$= \sum_{i=1}^{n} (\sigma^2 + \mu^2) - n \left(\frac{\sigma^2}{n} + \mu^2 \right)$$

$$= n\sigma^2 + n\mu^2 - \sigma^2 - n\mu^2$$

$$= (n-1)\sigma^2$$

Remember we have to mult by in 1/n at the end

Sample variance estimators

Simplifying the expression in the expectation:

$$= \left(\sum_{i=1}^{n} \mathbb{E}[X_i^2] \right) - \mathbb{E}[n\mu^2]$$

$$= n\sigma^2 + n\mu^2 - \sigma^2 - n\mu^2$$

$$= (n-1)\sigma^2$$

Remember we have to mult by in 1/n at the end

$$\text{bias}(\hat{\sigma^2}_n) = \frac{n-1}{n}\sigma^2 - \sigma^2$$

Sample variance estimators

$$\hat{\sigma^2}_n = \frac{1}{n}\sum_{i=1}^{n}(X_i - \mu)^2$$

By plugging in our estimator into the definition of bias, making the i.i.d. assumption and being comfortable with manipulating expectation identities, we showed the bias is non-zero:

$$\text{bias}(\hat{\sigma^2}_n) = \frac{n-1}{n}\sigma^2 - \sigma^2$$

Variance of an estimator

Definition 2.1. Statistic A statistic S is a random variable that is a function of some data $\mathcal{D}$, $S = g(\mathcal{D})$ where the data $\mathcal{D}$ is a collection of random variables.

An estimator is a statistic that aims at approximating a parameter/property of a distribution.

$$\mathrm{Var}(\hat{S}) = \mathbb{E}\left[(\hat{S} - \mathbb{E}(\hat{S}))^2\right]$$

Visualization of bias-variance

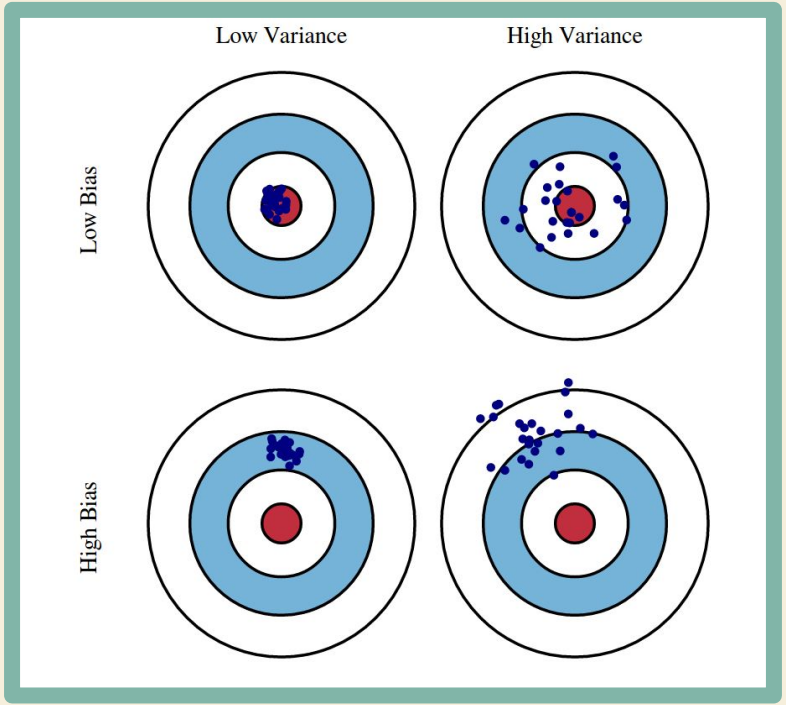

Bias-variance decomposition

$$\mathrm{Err}(\hat{S}) = \mathbb{E}[(\hat{S} - S)^2] = \mathrm{Bias}^2 + \mathrm{Var}$$

error of $\hat{S}$ is Bias^2 + variance $E[\hat{S}-S]$ + $E[(\hat{S}-E(\hat{S}))^2]$

The bias-variance decomposition simply tells us that the (mean squared error) of an estimator is completely accounted for by the sum of the squared bias and variance of the estimator.

Bias-variance decomposition

$$\mathrm{Err}(\hat{S}) = \mathbb{E}[(\hat{S} - S)^2] = \mathrm{Bias}^2 + \mathrm{Var}$$

$$= \mathbb{E}[(g(\mathcal{D}) - S)^2]$$

only works for MSE of estimator

The bias-variance decomposition simply tells us that the *(mean squared error)* of an estimator is completely accounted for by the sum of the squared bias and variance of the estimator.

bias: $E[\hat{S} - S]$ var: $E[(E(\hat{S}) - \hat{S})^2]$

Bias-variance proof

$$\mathbb{E}[(\hat{S} - S)^2] = \mathbb{E}\left[\left(\hat{S} - \mu + \mu - S\right)^2\right]$$

$$= \mathbb{E}\left[\left((\hat{S} - \mu) + (\mu - S)\right)^2\right]$$

$\pm \mu$

$= E\left[(\hat{S} - E[\hat{S}])^2\right] + E\left[(E[\hat{S}] - S)^2\right]$

Letting: $\mu = \mathbb{E}[\hat{S}]$

Bias-variance proof

$$
\begin{aligned}
\mathbb{E}[(\hat{S} - S)^2] &= \mathbb{E}\left[\left(\hat{S} - \mu + \mu - S\right)^2\right] \\
&= \mathbb{E}\left[\left((\hat{S} - \mu) + (\mu - S)\right)^2\right] \\
&= \mathbb{E}\left[(\hat{S} - \mu)^2 + 2(\hat{S} - \mu)(\mu - S) + (\mu - S)^2\right]
\end{aligned}
$$

Bias-variance proof

$$\mathbb{E}[(\hat{S} - S)^2] = \mathbb{E}\left[\left(\hat{S} - \mu + \mu - S\right)^2\right]$$

$$= \mathbb{E}\left[\left((\hat{S} - \mu) + (\mu - S)\right)^2\right] \qquad \mu = E[\hat{S}]$$

$$= \mathbb{E}\left[(\hat{S} - \mu)^2 + 2(\hat{S} - \mu)(\mu - S) + (\mu - S)^2\right]$$

$$= \mathbb{E}\left[(\hat{S} - \mu)^2 + (\mu - S)^2\right] \qquad (\star)$$

$E[E[\hat{S}]] = \hat{S}$

$E\left[(\hat{S} - \mu)(\mu - S)\right] = E\left[\hat{S}\mu - \hat{S}S - \mu\mu + \mu S\right] = E[\hat{S}E[S]] - E[\hat{S}S] - E[\mu^2] + E[\mu S]$

$= E[S]E[\hat{S}] - E[S]E[\hat{S}] - \mu^2 + \mu E[S]$

$= 0$

Bias-variance proof

$$\mathbb{E}[(\hat{S} - S)^2] = \mathbb{E}\left[\left(\hat{S} - \mu + \mu - S\right)^2\right]$$

$$= \mathbb{E}\left[\left((\hat{S} - \mu) + (\mu - S)\right)^2\right]$$

$$= \mathbb{E}\left[(\hat{S} - \mu)^2 + 2(\hat{S} - \mu)(\mu - S) + (\mu - S)^2\right]$$

$$= \mathbb{E}\left[(\hat{S} - \mu)^2 + (\mu - S)^2\right] \qquad (\star)$$

$$= \mathbb{E}\left[(\hat{S} - \overset{E(S)}{\mu})^2\right] + \mathbb{E}\left[(\overset{E(S)}{\mu} - S)^2\right]$$

$$= \mathbb{E}\left[(\hat{S} - \mu)^2\right] + (\mu - S)^2 = \text{var} + \text{bias}^2$$

bias $E(\hat{S} - S)$

$= E[\hat{S}] - E(S)$

$= \mu - S$

Bias-Variance in ML

Our course notes discuss the particulars of bias-variance decomposition in linear models, which you should study (in order to answer exercises), but in lecture we will be more general/conceptual.

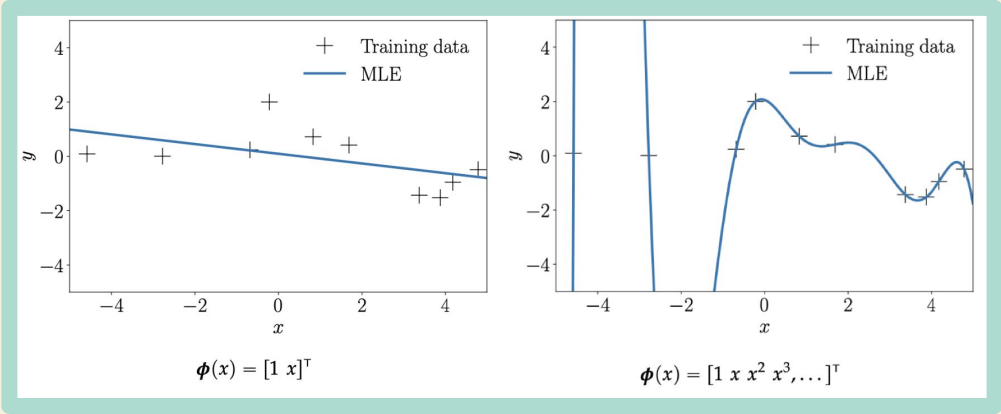

https://www.cs.cornell.edu/courses/cs4780/2022sp/notes/LectureNotes17.html

Bias-Variance in ML

$$\hat{y} = \int_y P(y|\mathbf{x})dy$$

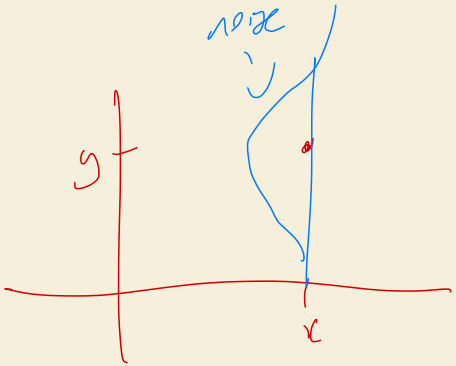

Consider the true label of our machine learning models to be the integral over the likelihood

Bias-Variance in ML

$$\hat{y} = \int_y y * P(y|\mathbf{x}) dy$$

We can write the error of our model as the integral over both the output and input space of our mean squared error:

$$\mathbb{E}_{P(\mathbf{x},y)}[(f^{\theta}(\mathbf{x}) - y)^2] = \int_x \int_y (f^{\theta}(\mathbf{x}) - y)^2 P(\mathbf{x}, y) dy d\mathbf{x}$$

MSE error

model

Bias-Variance in ML

$$\mathbb{E}_{P(\mathbf{x},y)}[(f^{\theta}(\mathbf{x}) - y)^2] = \int_x \int_y (f^{\theta}(\mathbf{x}) - y)^2 P(\mathbf{x}, y) dy d\mathbf{x}$$

$\hat{S}$ $-S$

need $E[\hat{S}]$

When we compute the bias and variance of an estimator (in this case the estimator of our risk of a model), we need to do so with respect to the "average model"

$$\hat{f}^{\theta} = \mathbb{E}_{\mathcal{D} \sim P(\mathbf{x},y)} \int_{\mathcal{D}} f^{\theta}_{\mathcal{D}} P(\mathcal{D}) d\mathcal{D}$$

avg. model given datasets

integrate over all potential datasets

Bias-Variance in ML

Now, we have defined all of the components we need to quantify the bias and variance of our model error estimator:

$$\mathbb{E}_{\mathbf{x},y,\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x}) - y)^2\right]$$

MSE

$E[\hat{S} - S] \rightarrow E[(\hat{S} - E(\hat{S}))^2] + E[(E(\hat{S}) - S)^2]$

Var + Bias

$\rightarrow E[(f_D^{\theta}(x) - \hat{f}^{\theta}(x))^2] + E[(\hat{f}^{\theta}(x) - y)^2]$

Bias-Variance in ML

$$\mathbb{E}_{\mathbf{x},y,\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x})-y)^2\right]$$

As before, we consider subtracting and adding the mean from the estimated quantity:

$$=\mathbb{E}_{\mathbf{x},y,\mathcal{D}}\left[\left((f_{\mathcal{D}}^{\theta}(\mathbf{x})-\hat{f}^{\theta}(\mathbf{x}))+(\hat{f}^{\theta}(\mathbf{x})-y)\right)^2\right]$$

Bias-Variance in ML

By the same cancellation logic, we get that the middle terms of completing the square drop off, so:

$$= \mathbb{E}_{\mathbf{x},y,\mathcal{D}}\left[\left((f_{\mathcal{D}}^{\theta}(\mathbf{x}) - \hat{f}^{\theta}(\mathbf{x}))^2 + (\hat{f}^{\theta}(\mathbf{x}) - y)^2\right)\right]$$

And then by linearity of expectation we have the definition of bias-variance again:

$$\mathbb{E}_{\mathbf{x},\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x}) - \hat{f}^{\theta}(\mathbf{x}))^2\right] + \mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x}) - y)^2]$$

how much model deviates around avg model w/ diff datasets (variance of model)

risk of any classifier

Bias-Variance in ML

$$\mathbb{E}_{\mathbf{x},\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x}) - \hat{f}^{\theta}(\mathbf{x}))^2\right] + \underbrace{\mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x}) - y)^2]}$$

Notice! Another MSE term here, this time the risk of our average model. We can repeat our bias-variance steps to decompose this further

Bias-Variance in ML

Again, we can add and subtract the estimator for y to further decompose this term:

$$\mathbb{E}_{\mathbf{x},\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x})-\hat{f}^{\theta}(\mathbf{x}))^2\right]+\underbrace{\mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x})-y)^2]}$$

$$\mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x})-y)^2]=\mathbb{E}_{\mathbf{x},y}\left[(\hat{f}^{\theta}(\mathbf{x})-\hat{y}+\hat{y}-y)^2\right]$$

Bias-Variance in ML

Again, we can add and subtract the estimator for y to further decompose this term:

$$\mathbb{E}_{\mathbf{x},\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x}) - \hat{f}^{\theta}(\mathbf{x}))^2\right] + \mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x}) - y)^2]$$

$$\mathbb{E}_{\mathbf{x},y}[(\hat{f}^{\theta}(\mathbf{x}) - y)^2] = \mathbb{E}_{\mathbf{x},y}\left[(\hat{f}^{\theta}(\mathbf{x}) - \hat{y} + \hat{y} - y)^2\right]$$

$$= \mathbb{E}_{\mathbf{x},y}[(\hat{y} - y)^2] + \mathbb{E}_{\mathbf{x}}\ \ [(\hat{f}^{\theta}(\mathbf{x}) - \hat{y})^2]$$

noise

bias^2

Final bias-variance decomposition

$$\mathbb{E}_{\mathbf{x},y,\mathcal{D}}\left[(f_{\mathcal{D}}^{\theta}(\mathbf{x})-y)^2\right]$$

We decomposed the error into three terms:

$$\underbrace{\mathbb{E}_{\mathbf{x},y}[(\hat{y}-y)^2]}_{\text{Noise}}+\underbrace{\mathbb{E}_{\mathbf{x}}\ \ [(\hat{f}^{\theta}(\mathbf{x})-\hat{y})^2]}_{\text{Squared bias}}+\underbrace{\mathbb{E}_{\mathbf{x},\mathcal{D}}[(\hat{f}_{\mathcal{D}}^{\theta}(\mathbf{x})-\hat{f}^{\theta}(\mathbf{x}))^2]}_{\text{Variance}}$$

Visualizing this decomposition

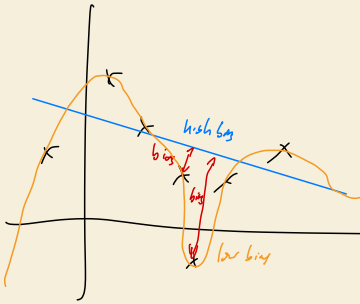

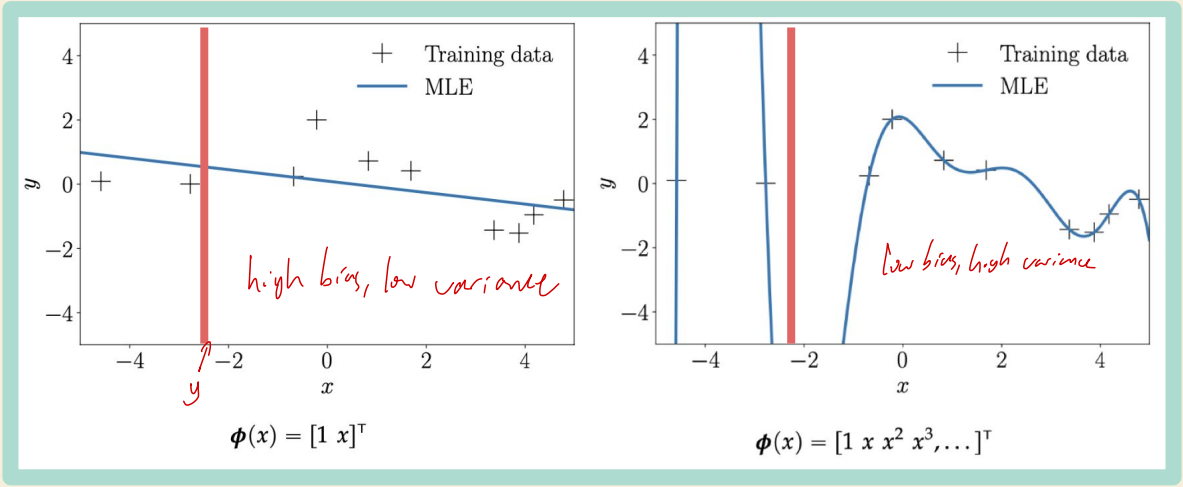

$$\underbrace{\mathbb{E}_{\mathbf{x},y}[(\hat{y} - y)^2]}_{\text{Noise}} + \underbrace{\mathbb{E}_{\mathbf{x}}\ [(\hat{f}^{\theta}(\mathbf{x}) - \hat{y})^2]}_{\text{Squared bias}} + \underbrace{\mathbb{E}_{\mathbf{x},\mathcal{D}}[(\hat{f}^{\theta}_{\mathcal{D}}(\mathbf{x}) - \hat{f}^{\theta}(\mathbf{x}))^2]}_{\text{Variance}}$$

Model selection

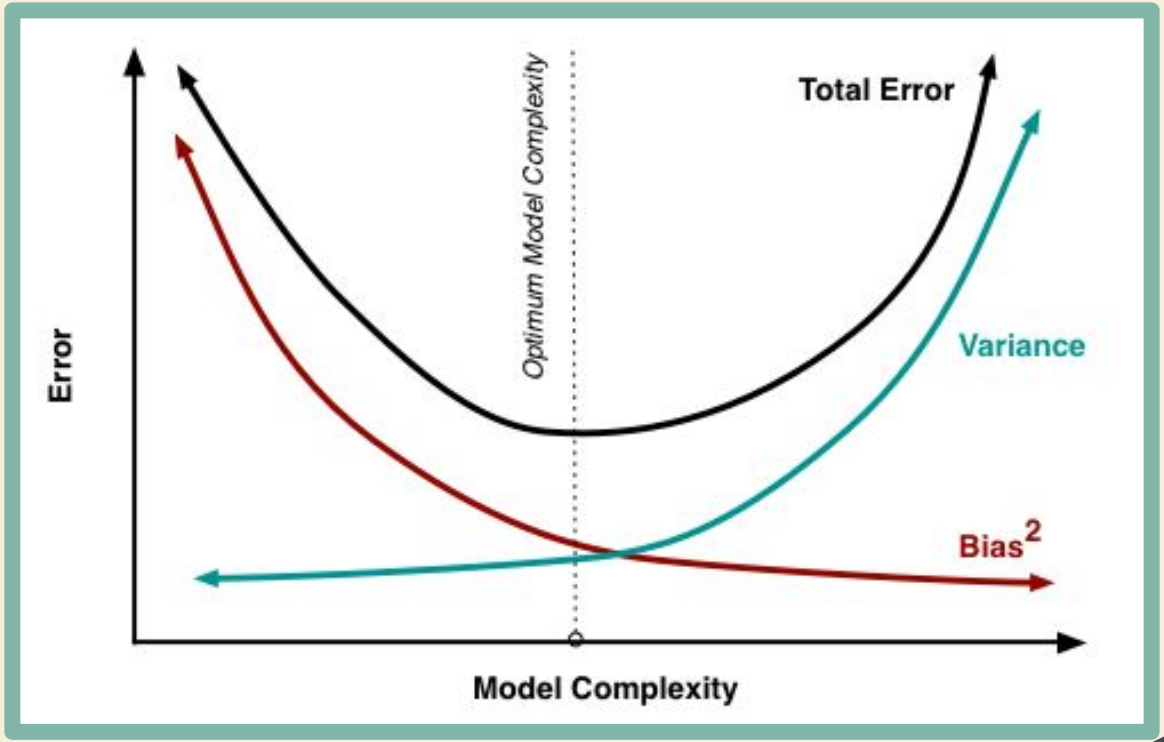

Next lecture:
Principal Component-Analysis